

Amo Samuel

Report on Transaction Management and Caching Implementation

This document provides an overview of the implementation of transaction management and caching mechanisms in the Spring Data project. It covers the use of transaction management for ensuring data consistency and the application of caching to improve performance.

Overview

Transaction management is crucial for maintaining data integrity and consistency, especially when multiple operations need to be executed as a single unit of work. In this project, I use Spring's transaction management capabilities to handle transactions explicitly.

Implementation

PatientService Class

In the **PatientService** class, I have implemented explicit transaction management for creating a new **Patient** record.

Transaction Definition: I define a transaction using **DefaultTransactionDefinition**. This includes setting the transaction name and propagation behavior.

```
DefaultTransactionDefinition def = new DefaultTransactionDefinition();
def.setName("Patient transaction");
def.setPropagationBehavior(TransactionDefinition.PROPGATION_REQUIRED);
TransactionStatus status = transactionManager.getTransaction(def);
```

Transaction Management: I obtain the transaction status and perform operations within a try-catch block. If all operations succeed, we commit the transaction; otherwise, I roll back the transaction to ensure data consistency.

```

        transactionManager.commit(status);
        return patientRepository.save(patient);

    }catch (Exception exception) {
        transactionManager.rollback(status);
        throw exception;
    }
}

```

Caching

Caching is used to store frequently accessed data in memory to improve performance and reduce database load. In this project, I use Caffeine as the caching solution.

Configuration

Caching is configured in the `application.properties` file. I specified Caffeine as the cache provider and configured cache properties such as expiration and maximum size.

```

4
5 # cache configuration
6 spring.cache.type=caffeine
7 spring.cache.cache-names=patients
8 spring.cache.caffeine.spec=expireAfterWrite=60m,maximumSize=100
9
10
11

```

expireAfterWrite=60m: Specifies that cache entries will expire 60 minutes after being written.

maximumSize=100: Limits the cache to 100 entries to prevent excessive memory usage.

PatientService Class

Caching Implementation:

In the `PatientService` class, caching is applied to methods to reduce database calls:

Caching All Patients: The `getAllPatients()` method is annotated with `@Cacheable`, which caches the result of the method call.

```

5      }
6
7      @Cacheable("patients") 1 usage  samuelamo001
8      public List<Patient> getAllPatients() {
9          return patientRepository.findAll();
10     }
11
12

```

Evicting Cache on Update: The `updatePatient()` method is annotated with `@CacheEvict`, which clears the cache entries when a patient record is updated.

```

12
13  @CacheEvict(value = "patients", allEntries = true) 1 usage  samuelamo001
14  @
15  public Patient updatePatient(Long patientId, PatientDTO patientDTO) {
16
17      Patient patient = patientRepository.findById(patientId).orElseThrow(() -> new IT
18      patient.setName(patientDTO.getName());
19      patient.setSurname(patientDTO.getSurname());
20      patient.setAddress(patientDTO.getAddress());
21      patient.setTelephone(patientDTO.getTelephone());
22      patient.setDiagnosis(patientDTO.getDiagnosis());
23      patient.setBedNumber(patientDTO.getBedNumber());
24
25      return patientRepository.save(patient);
26  }
27

```

Conclusion

The combination of transaction management and caching in this Spring Data project enhances data consistency, integrity, and performance. By explicitly managing transactions and leveraging Caffeine for caching, I ensured that the application operates efficiently and reliably.